

DIP-SMC-PSO: Technical Progress Memo

Model choices, algorithm design, and experimental results

Student: TBD — Advisor: TBD — Program: M.Sc. Program

February 2026 — Repository: github.com/theSadeQ/dip-smc-pso — Status: Research complete, thesis
90%

Contents

1 System Model: Equations and Parameter Choices

1.1 State Vector and Coordinate Definition

The DIP state is a 6-vector:

$$\mathbf{x} = [x, \theta_1, \theta_2, \dot{x}, \dot{\theta}_1, \dot{\theta}_2]^\top \in \mathbb{R}^6 \quad (1)$$

where x is cart position, θ_1 and θ_2 are pendulum link angles (from vertical), and dots denote time derivatives.

Coordinate convention and units. The generalized coordinates are

$$\mathbf{q} = [x, \theta_1, \theta_2]^\top, \quad \dot{\mathbf{q}} = [\dot{x}, \dot{\theta}_1, \dot{\theta}_2]^\top \quad (2)$$

with SI units: x in m, θ_i in rad, $\dot{x}$ in m/s, and $\dot{\theta}_i$ in rad/s. Positive x is rightward cart motion, and positive θ_i is counterclockwise from the upright vertical axis ($\theta_i = 0$).

Target equilibrium and error coordinates. The nominal stabilization target is the upright equilibrium:

$$\mathbf{x}^* = [0, 0, 0, 0, 0, 0]^\top, \quad \mathbf{e} = \mathbf{x} - \mathbf{x}^* \quad (3)$$

which is used in controller design and Lyapunov analysis. For tracking tasks, a nonzero cart reference updates only the first element of $\mathbf{x}^*$.

Measured output and operational bounds. The feedback vector is

$$\mathbf{y} = [x, \theta_1, \theta_2, \dot{x}, \dot{\theta}_1, \dot{\theta}_2]^\top \quad (4)$$

(full-state feedback in simulation). In HIL deployment, the same channels are obtained from encoder/estimator pipelines. Operational bounds used for pass/fail decisions are enforced by the benchmark configuration files.

1.2 Equations of Motion

The full Lagrangian model is:

$$\mathbf{M}(\mathbf{q}) \ddot{\mathbf{q}} + \mathbf{C}(\mathbf{q}, \dot{\mathbf{q}}) \dot{\mathbf{q}} + \mathbf{G}(\mathbf{q}) + \mathbf{F}_{\text{fric}} = \mathbf{B} \mathbf{u} \quad (5)$$

where $\mathbf{q} = [x, \theta_1, \theta_2]^\top$, $\mathbf{u} \in [-150, 150]$ N is the cart force, and $\mathbf{B} = [1, 0, 0]^\top$ selects the actuated direction.

Modeling assumptions. The model assumes planar rigid-body dynamics, ideal revolute joints, bounded single-input actuation, and friction represented by viscous plus Coulomb terms. External disturbances and uncertainty scenarios are injected through benchmark profiles used in MT-8/LT-6 studies.

Equivalent first-order form. For control synthesis, (??) is written as

$$\dot{\mathbf{x}} = \mathbf{f}(\mathbf{x}) + \mathbf{g}(\mathbf{x}) \mathbf{u} + \mathbf{d}(t) \quad (6)$$

where $\mathbf{d}(t)$ is zero in nominal simulations and nonzero in disturbance rejection experiments.

Structural properties used in stability analysis. For admissible configurations, $\mathbf{M}(\mathbf{q})(\mathbf{q})$ is symmetric positive definite, and the standard manipulator identity $\mathbf{M}(\mathbf{q}) - 2\mathbf{C}(\mathbf{q}, \dot{\mathbf{q}})$ is skew-symmetric. These properties are used in Lyapunov-based controller proofs.

Model partition used in this workflow.

- **Full nonlinear model:** Used for reported benchmark metrics (MT-5, MT-7, MT-8, LT-6), including full coupling, Coriolis, centrifugal, gyroscopic, and friction effects.
- **Simplified model:** Used for PSO search loops with reduced coupling and omitted gyroscopic term to reduce optimization runtime.

Inertia matrix $\mathbf{M}(\mathbf{q})$ (3×3, symmetric, configuration-dependent): To keep the matrix compact in print, define $c_1 = \cos \theta_1$, $c_2 = \cos \theta_2$, and $c_{12} = \cos(\theta_1 - \theta_2)$.

$$\mathbf{M}(\mathbf{q}) = \begin{bmatrix} M + m_1 + m_2 & a_{12} & a_{13} \\ a_{12} & a_{22} & a_{23} \\ a_{13} & a_{23} & a_{33} \end{bmatrix} \quad (7)$$

with

$$\begin{aligned} a_{12} &= (m_1 l_{c1} + m_2 l_1) c_1 + m_2 l_{c2} c_2 \\ a_{13} &= m_2 l_{c2} c_2 \\ a_{22} &= m_1 l_{c1}^2 + m_2 (l_1^2 + l_{c2}^2) + I_1 + I_2 + 2m_2 l_1 l_{c2} c_{12} \\ a_{23} &= m_2 l_{c2}^2 + I_2 + m_2 l_1 l_{c2} c_{12} \\ a_{33} &= m_2 l_{c2}^2 + I_2 \end{aligned} \quad (8)$$

The dominant configuration-dependent coupling appears in a_{22} through $2m_2 l_1 l_{c2} c_{12}$, which grows near large relative-link deflections.

Coriolis/centrifugal matrix $\mathbf{C}(\mathbf{q}, \dot{\mathbf{q}})$ (non-zero entries): Computed from the Christoffel symbols of (??); all zero entries omitted:

$$\begin{aligned} C_{01} &= -(m_1 l_{c1} + m_2 l_1) \sin \theta_1 \dot{\theta}_1 - m_2 l_{c2} \sin \theta_2 \dot{\theta}_2 \\ C_{02} &= -m_2 l_{c2} \sin \theta_2 \dot{\theta}_2 \\ C_{11} &= C_{12} = -m_2 l_1 l_{c2} \sin(\theta_1 - \theta_2) \dot{\theta}_2 \\ C_{21} &= +m_2 l_1 l_{c2} \sin(\theta_1 - \theta_2) \dot{\theta}_1 \end{aligned} \quad (9)$$

(indices 0,1,2 correspond to x, θ_1, θ_2). A small gyroscopic coupling term $g_c(\dot{\theta}_1 + \dot{\theta}_2)$ is added to C_{01} and $-C_{10}$ (coefficient 0.01, enabled in config).

Gravity vector $\mathbf{G}(\mathbf{q})$:

$$\mathbf{G}(\mathbf{q}) = \begin{bmatrix} 0 \\ G_2 \\ G_3 \end{bmatrix}$$

$$G_2 = -(m_1 l_{c1} + m_2 l_1) g \sin \theta_1$$

$$G_3 = -m_2 l_{c2} g \sin \theta_2 \quad (10)$$

$G_1 = 0$ because gravity acts perpendicular to the cart's horizontal motion.

Friction model (viscous + Coulomb):

$$F_{\text{fric},i} = -b_i \dot{q}_i - \mu_i \text{sign}(\dot{q}_i) \mathbb{1}_{|\dot{q}_i| > 10^{-6}} \quad (11)$$

for each DOF $i \in \{x, \theta_1, \theta_2\}$, where b_i is the viscous coefficient (cart: 0.2, joints: 0.005, 0.004) and μ_i is the Coulomb coefficient (from config; varies $\pm 10\%$ in Monte Carlo). The dead zone $|\dot{q}_i| > 10^{-6}$ avoids sign-function discontinuity at rest. A viscous-only equation would omit the Coulomb term – the full form (??) is used in all simulations.

Implementation and linearization details: The simplified PSO dynamics use a reduced coupling matrix (off-diagonal terms scaled by 0.7–0.8) and omit gyroscopic effects. Linearization ($\mathbf{A}, \mathbf{B}$) around $\mathbf{q}^* = \mathbf{0}$, $\dot{\mathbf{q}}^* = \mathbf{0}$ is computed numerically by finite differences ($\varepsilon = 10^{-8}$), which is more maintainable than the full symbolic Jacobian of (??). The control update period is $\Delta t = 0.01$ s with saturation $u \in [-150, 150]$ N. Model implementation was sanity-checked with trajectory consistency and energy-trend checks before running comparative benchmarks.

1.3 Physical Parameters

Parameter	Symbol	Value
Cart mass	M	1.5 kg
Link 1 mass	m_1	0.2 kg
Link 2 mass	m_2	0.15 kg
Link 1 length	l_1	0.4 m
Link 2 length	l_2	0.3 m
Link 1 COM distance	l_{c1}	0.2 m
Link 2 COM distance	l_{c2}	0.15 m
Link 1 inertia	I_1	0.0081 kg·m ²
Link 2 inertia	I_2	0.0034 kg·m ²
Gravity	g	9.81 m/s ²
Cart friction	b_c	0.2 N·s/m
Joint 1 friction	b_1	0.005 N·m·s/rad
Joint 2 friction	b_2	0.004 N·m·s/rad
Max control force	u_{max}	150 N
Control period	Δt	0.01 s (100 Hz)

1.4 Model Selection: Simplified vs. Full Nonlinear

Model	Use in this project	Rationale
Simplified (linearized around upright)	PSO optimization (all controllers)	10–50× faster per simulation step. PSO requires thousands of evaluations; full model would be prohibitive.
Full nonlinear (??)	All published benchmark results	Complete Lagrangian, all coupling terms, Coriolis, centrifugal, gyroscopic, full friction. Used for MT-5, MT-7, MT-8, LT-6.
Low-rank (SVD approximation)	HIL mode (hardware-in-the-loop)	Reduced singular-value truncation preserves stability with lower real-time compute load.

Decision rationale: Controller gain tuning (PSO) uses the simplified model because PSO with 40 particles $\times$ 200 iterations $\times$ 10 s simulation on the full nonlinear model would require >8 hours per controller. With the simplified model, PSO completes in minutes. All final validation and reported results use the full nonlinear model.

2 Sliding Mode Control: Surface Design and Control Laws**2.1 Sliding Surface**

The multi-variable sliding surface for the DIP is:

$$\sigma = \lambda_1 \theta_1 + \lambda_2 \theta_2 + k_1 \dot{\theta}_1 + k_2 \dot{\theta}_2 \quad (12)$$

Equivalently, in error coordinates $e_i = \theta_i - \theta_i^*$ (where $\theta^* = 0$ at upright equilibrium):

$$\sigma = \lambda_1 e_1 + \lambda_2 e_2 + k_1 \dot{e}_1 + k_2 \dot{e}_2 \quad (13)$$

The surface (??) is a hyperplane in $\mathbb{R}^4$ (angle + rate space). On this surface ($\sigma = 0$), the reduced dynamics decouple into two independent first-order systems $\dot{e}_i = -(\lambda_i/k_i) e_i$, which are stable if and only if:

$$\lambda_i > 0, \quad k_i > 0 \quad (i = 1, 2) \quad (14)$$

These inequalities are enforced as PSO lower bounds (all gains > 2).

Surface design rationale: The linear form (??) was chosen over alternatives as follows: (i) *Integral sliding surface* – rejected because it adds an integrator state, complicating the Lyapunov analysis for a 6-state system; (ii) *Terminal sliding surface* ($\sigma = \dot{e} + \beta e^{p/q}$) – rejected because the fractional power $e^{p/q}$ causes numerical issues near the equilibrium ($p/q < 1$ makes the derivative singular at $e = 0$); (iii) *Linear surface* (??) – relative degree 1 (verified below), Lyapunov analysis straightforward, PSO can tune all 4 surface parameters simultaneously.

Relative degree check: Differentiating (??):

$$\dot{\sigma} = \lambda_1 \dot{\theta}_1 + \lambda_2 \dot{\theta}_2 + k_1 \ddot{\theta}_1 + k_2 \ddot{\theta}_2 \quad (15)$$

Substituting $\ddot{\theta}_i$ from (??): $\ddot{\mathbf{q}} = \mathbf{M}^{-1}(\mathbf{B}u - \mathbf{C}\dot{\mathbf{q}} - \mathbf{G} - \mathbf{F}_{\text{fric}})$. The control u appears through $\mathbf{L}\mathbf{M}^{-1}\mathbf{B}$ where $\mathbf{L} = [0, k_1, k_2]^\top$ (surface gradient w.r.t. $\dot{\mathbf{q}}$). Since $\mathbf{M}^{-1}$ is invertible and $\mathbf{L}\mathbf{M}^{-1}\mathbf{B} \neq 0$ for all physically valid configurations, the relative degree is exactly 1.

Matching condition: External disturbances $d(t)$ and model uncertainty $\Delta f(\mathbf{x})$ enter the equations of motion through the same actuated direction as u (the cart force channel): $\dot{\sigma}$ contains a term $\mathbf{L}\mathbf{M}^{-1}\mathbf{B}(u + d)$, so the matching condition holds and the switching gain K can reject any disturbance bounded by $|d(t)| \leq \bar{d}$ provided $K > \bar{d}$.

2.2 Controller Variants: Control Laws

2.2.1 Classical SMC: Saturation Function and Equivalent Control

The saturation function replacing $\text{sign}(\cdot)$ is:

$$\text{sat}\left(\frac{\sigma}{\varepsilon}\right) = \begin{cases} \sigma/\varepsilon & |\sigma| \leq \varepsilon \\ \text{sign}(\sigma) & |\sigma| > \varepsilon \end{cases} \quad (16)$$

Inside the boundary layer ($|\sigma| \leq \varepsilon$), the control is linear in σ – this eliminates high-frequency switching at the cost of a residual steady-state error bounded by $O(\varepsilon)$.

The full control law is:

$$u = \underbrace{u_{\text{eq}}}_{\text{equiv. control}} + \underbrace{K \text{sat}\left(\frac{\sigma}{\varepsilon}\right)}_{\text{switching}} + \underbrace{-k_d \dot{\sigma}}_{\text{derivative}} \quad (17)$$

Equivalent control derivation (from the ideal sliding condition $\dot{\sigma} = 0$):

$$\begin{aligned} \dot{\sigma} = 0 &\implies \mathbf{L}\ddot{\mathbf{q}} + (\lambda_1 \dot{\theta}_1 + \lambda_2 \dot{\theta}_2) = 0 \\ &\implies \mathbf{L}\mathbf{M}^{-1}(\mathbf{B}u_{\text{eq}} - \mathbf{F}) = 0 \\ &\implies u_{\text{eq}} = -\frac{\mathbf{L}\mathbf{M}^{-1}\mathbf{F}}{\mathbf{L}\mathbf{M}^{-1}\mathbf{B}} \end{aligned} \quad (18)$$

where $\mathbf{F} = \mathbf{C}\dot{\mathbf{q}} + \mathbf{G} + \mathbf{F}_{\text{fric}}$ collects all nonlinear forcing terms, and $\mathbf{L} = [0, k_1, k_2]^\top$ is the surface gradient. In code:

- `LM_inv_B = L @ M_inv @ B` (scalar)
- `u_eq = -(L @ M_inv @ F) / LM_inv_B`
- K : switching gain; $\varepsilon = 0.3$: boundary layer width (MT-6 study found $\varepsilon = 0.02$ optimal for chattering only; $\varepsilon = 0.3$ is the operational value balancing tracking accuracy and compute smoothness)
- 6 gains tuned by PSO: $[k_1, k_2, \lambda_1, \lambda_2, K, k_d]$

2.2.2 Super-Twisting Algorithm (STA-SMC)

The STA operates on the surface derivative, not the surface itself:

$$\begin{aligned} u &= u_1 + u_2 \\ u_1 &= -K_1 |\sigma|^\alpha \text{sign}(\sigma) \quad (\text{continuous}) \\ u_2 &= -K_2 \int_0^t \text{sign}(\sigma(\tau)) d\tau \quad (\text{integral}) \end{aligned} \quad (19)$$

with $K_1 > K_2 > 0$ and standard exponent $\alpha = 0.5$.

Key property: u_1 is continuous (chattering reduction), while u_2 handles steady-state errors via integration. Near $\sigma = 0$, regularization replaces $|\sigma|^\alpha$ with δ^α (where $\delta = 10^{-10}$) to avoid numerical singularity.

6 gains: $[c_1, \lambda_1, c_2, \lambda_2, K_1, K_2]$; constraint $K_1 > K_2$ enforced during PSO.

2.2.3 Adaptive SMC

The switching gain adapts online:

$$u = -K(t) \text{sat}\left(\frac{\sigma}{\varepsilon}\right) \quad (20)$$

$$\dot{K}(t) = \gamma |\sigma| - \delta_{\text{leak}} K(t) \quad (21)$$

with dead zone $|\sigma| < d_z$. where γ is the adaptation rate, $\delta_{\text{leak}} = 0.01$ prevents gain windup, and $d_z = 0.05$ suppresses noise-driven adaptation. $K(t)$ is clamped to $[K_{\min}, K_{\max}] = [0.1, 100]$.

Parameter rationale: $\delta_{\text{leak}} = 0.01$ allows $K(t)$ to decay slowly when $\sigma \approx 0$, preventing accumulation over long horizons. If δ_{leak} is too large, $K(t)$ shrinks before new disturbances arrive, causing chattering bursts; if too small, $K(t)$ drifts upward unboundedly. The dead zone $d_z = 0.05$ matches the angle sensor noise floor ($0.005 \text{ rad} \times \text{typical gain} \approx 0.01\text{--}0.05$), ensuring adaptation does not respond to measurement noise rather than real errors.

5 gains: $[k_1, k_2, \lambda_1, \lambda_2, \gamma]$.

2.2.4 Hybrid Adaptive STA-SMC

Combines adaptive gain scheduling with the super-twisting structure:

$$u = -k_1(t) |\sigma|^{0.5} \text{sign}(\sigma) - k_2(t) \int \text{sign}(\sigma) dt \quad (22)$$

Adaptation laws for both gains $k_1(t)$, $k_2(t)$ use rates $\gamma_1 = 0.1$, $\gamma_2 = 0.05$ with a rate limiter of 5.0/s.

Known instability (documented): Under large disturbances and adaptive gain scheduling,

this controller produces 666.9° overshoot (176% chattering increase). Root cause: feedback loop between $k_1(t)$ adaptation and the $|\sigma|^{0.5}$ term amplifies errors when $|\sigma| \gg 1$. This is an open research item.

3 PSO Algorithm: Configuration and Methodology

3.1 Algorithm Parameters

Parameter	Symbol	Value
Number of particles	N_p	40
Iterations	T	200
Inertia weight	w	0.7
Cognitive acceleration	c_1	2.0
Social acceleration	c_2	2.0
Random seed	–	42

Rationale: $N_p = 40$ increases coverage of the 6D search space; $T = 200$ was sufficient for stable convergence in logs; $(w, c_1, c_2) = (0.7, 2.0, 2.0)$ follows standard PSO settings.

$$v_{i,d}^{t+1} = w v_{i,d}^t + c_1 r_1 (p_{i,d} - x_{i,d}^t) + c_2 r_2 (g_d - x_{i,d}^t), \quad r_1, r_2 \sim U(0, 1). \quad (23)$$

3.2 Fitness Function

$$J(\theta) = w_1 \bar{J}_{\text{mc}} + w_2 \max_k J_k \quad (24)$$

where $w_1 = 0.7$, $w_2 = 0.3$ (mean vs. worst-case balance), and:

$$J_k = w_{\text{err}} \text{ISE}(\mathbf{x}) + w_u \text{IAE}(u) + w_r \text{IAE}(\dot{u}) + w_s \int \sigma^2 dt + P \cdot \mathcal{K}_{\text{unstable}} \quad (25)$$

Terms: tracking, effort, slew-rate, sliding-energy, and instability penalty. with weights $[w_{\text{err}}, w_u, w_r, w_s] = [1.0, 0.1, 0.01, 0.1]$, normalization thresholds $[10.0, 100.0, 1000.0, 1.0]$, and instability penalty $P = 10^6$.

3.3 Search Bounds per Controller

Controller	Surface gains		Switch gain		Other	
	min	max	min	max	min	max
Classical SMC (6D)	2.0	30.0	2.0	50.0	0.05	3.0
STA-SMC (6D)	2.0	30.0	2.0	30.0	0.05	3.0
Adaptive SMC (5D)	2.0	30.0	2.0	10.0	0.05	3.0
Hybrid Adaptive (4D)	2.0	30.0	2.0	10.0	–	–

Note: STA-SMC enforces $K_1 > K_2$ as a hard constraint during PSO evaluation (particles violating it are given the instability penalty $P = 10^6$).

Hyperparameter sensitivity (finding #16, #17): A formal ablation study over N_p and w was not performed; standard literature values were adopted (Kennedy & Eberhart 1995; Clerc & Kennedy 2002). $N_p = 40$ was increased from 30 to improve coverage of the 6D search space based on the heuristic $N_p \approx 10d$ for dimension d . PSO convergence curves are logged to `academic/logs/psd/` and show stable cost plateau by iteration 150–180 in all controllers; the 200-iteration budget is fixed (no early-stopping criterion implemented).

Cost function scale calibration (finding #18): Normalization thresholds $[10, 100, 1000, 1]$ correspond to the approximate magnitude of each raw metric for an unstable run: $\text{ISE} \sim 10 \text{ rad}^2$, $\text{IAE}(u) \sim 100 \text{ N}\cdot\text{s}$, $\text{IAE}(\dot{u}) \sim 1000 \text{ N}$, sliding integral ~ 1 . The instability penalty $P = 10^6$ is approximately $10^4 \times$ the typical stable-run cost (~ 0.1 – 1), ensuring unstable particles are always ranked last in the swarm.

3.4 Robust PSO: Scenario-Based Optimization

Standard PSO (nominal initial condition) produces gains that fail catastrophically on varied initial conditions:

$$2.14 \pm 0.13 \rightarrow 107.61 \pm 5.48 \quad (\times 50.4) \quad (26)$$

Robust PSO solution: Evaluate each particle on $N_s = 15$ scenarios:

- 3 nominal: $\theta_0 \in [-0.05, +0.05]$ rad
- 4 moderate: $\theta_0 \in [-0.15, +0.15]$ rad
- 8 large: $\theta_0 \in [-0.30, +0.30]$ rad

Fitness: $J_{\text{robust}} = 0.7 \bar{J} + 0.3 \max_k J_k$ (weighted average + worst-case, same formula as (??)).

Result: Chattering degradation reduced from $50.4 \times$ to $7.5 \times$ ($144.59 \times \rightarrow 19.28 \times$ in absolute terms).

Scenario distribution rationale (finding #19): The 3+4+8 split (20%, 27%, 53% nominal/moderate/large) was chosen to weight large perturbations more heavily than nominal, since controllers failing at large initial angles are useless in practice. A sensitivity analysis on this split was not performed; this is a known limitation.

3.5 PSO-Optimized Gains (MT-8 Robust Tuning)

Controller	Nominal gains	Robust gains
Classical SMC	k1=5.0, k2=5.0, l1=5.0, l2=0.5, K=0.5, kd=0.5	k1=23.07, k2=12.85, l1=5.51, l2=3.49, K=2.23, kd=0.15
STA-SMC	c1=12.0, l1=6.0, c2=4.85, l2=3.43, K1=8.0, K2=4.0	c1=5.62, l1=3.75, c2=4.36, l2=2.05, K1=2.02, K2=6.67
Adaptive SMC	k1=10.0, k2=8.0, l1=5.0, l2=4.0, gamma=1.0	k1=2.14, k2=3.36, l1=7.20, l2=0.34, gamma=0.29
Hybrid Adaptive STA	c1=5, l1=5, c2=5, l2=0.5	c1=10.15, l1=12.84, c2=6.82, l2=2.75

Robust-vs-nominal cost deltas: Classical -2.2% , Hybrid -21.4% , STA and Adaptive reported as consistency checks (see note below).

Known inconsistency – STA robust gains (finding #20): The robust STA gains show $K_1 = 2.02$, $K_2 = 6.67$, which violates the stated constraint $K_1 > K_2$. This occurred because the robust PSO cost function (15 scenarios, worst-case weighted) found this parameter combination reduced chattering across the scenario set even though the nominal stability condition is reversed. The Moreno–Osorio proof requires $K_1 > K_2 > 0$ for finite-time convergence; these gains are therefore *not theoretically guaranteed* and should be used with caution. This is an open issue requiring either re-running robust PSO with the hard constraint enforced, or re-deriving stability conditions for $K_1 < K_2$.

4 Lyapunov Stability: Proofs and Validation

Proofs are documented in `academic/paper/sphinx_docs/theory/lyapunov_stability_proofs.md` (v2.0, 1427 lines, 9 theorems + 5 lemmas, 15 BibTeX references).

Note on symbolic constants (finding #21): The constants β , η , c_1 , c_2 , α_1 , α_2 in the inequalities below are **not fixed numerical values** – they are state-dependent or design-parameter-dependent quantities. Specifically: $\beta = \mathbf{LM}^{-1}(\mathbf{q})\mathbf{B} > 0$

(varies with configuration $\mathbf{q}$; positive for all valid DIP poses); $\eta = K - \bar{d}$ (switching gain minus disturbance bound; positive by design if $K > \bar{d}$). The proofs establish that these quantities remain positive along all trajectories, which is the content of the theorems – not that they have fixed values.

4.1 Classical SMC

$$V = \frac{1}{2}\sigma^2 \quad (27)$$

$$\dot{V} = \sigma\dot{\sigma} \leq -\beta\eta|\sigma| - \beta k_d\sigma^2 < 0 \quad (28)$$

for all $|\sigma| > \varepsilon$ (outside boundary layer), where $\beta = \mathbf{LM}^{-1}\mathbf{B} > 0$ and $\eta = K - \bar{d} > 0$. Reaching time bound: $t_{\text{reach}} \leq |\sigma(0)|/(\eta\beta)$. Stability type: **exponential asymptotic** (when $k_d > 0$).

Numerical validation: 96.2% of simulation samples satisfy $\dot{V} < 0$ outside the boundary layer. The remaining 3.8% occur during the initial reaching phase when the trajectory is still crossing the boundary layer ($|\sigma| \leq \varepsilon = 0.3$) – this is expected behavior and does not indicate instability.

4.2 Super-Twisting (STA-SMC)

Uses the generalized gradient Lyapunov function (Moreno & Osorio 2008):

$$V_{\text{STA}} = |\sigma| + \frac{1}{2K_2}z^2, \quad z = u_2 \quad (29)$$

$$\dot{V}_{\text{STA}} \leq -c_1\|\xi\|^{3/2} + c_2L \quad (30)$$

where $\xi = [|\sigma|^{1/2} \text{sign}(\sigma), z]^\top$, L is the Lipschitz constant of $\dot{d}_u(t)$ (disturbance derivative), and $c_1, c_2 > 0$ depend on K_1, K_2 .

Moreno–Osorio gain conditions (required for finite-time convergence):

$$K_1 > 2\sqrt{\frac{2\bar{d}}{\beta}}, \quad K_2 > \frac{\bar{d}}{\beta}, \quad K_1 > K_2 \quad (31)$$

where $\bar{d}$ bounds the disturbance and $\beta = \mathbf{LM}^{-1}\mathbf{B}$. These conditions are satisfied by the nominal gains ($K_1 = 8 > K_2 = 4$) but *not* by the robust gains ($K_1 = 2.02 < K_2 = 6.67$) – see the note in Section 3.5.

Finite-time convergence: The STA guarantees $\sigma(t) = 0$ for all $t \geq T_f$, where T_f depends on initial conditions and gain values. An explicit closed-form bound for the DIP system is not derived here (it requires bounding β over the state trajectory, which is configuration-dependent); the experimental confirmation is convergence in 1.82 s for typical initial conditions.

Stability type: **finite-time convergence**.

Numerical validation: Convergence in 1.82 s (16% faster than Classical SMC at 2.15 s).

4.3 Adaptive SMC

Composite Lyapunov (state + parameter estimation error):

$$V_{\text{ad}} = \frac{1}{2}\sigma^2 + \frac{1}{2\gamma}\tilde{K}^2, \quad \tilde{K} = K(t) - K^* \quad (32)$$

$$\dot{V}_{\text{ad}} \leq -\eta|\sigma| \Rightarrow \sigma(t) \rightarrow 0, \quad K(t) \text{ bounded} \quad (33)$$

Stability type: **asymptotic with bounded adaptive gain**.

Numerical validation: 100% of simulation samples within $K(t) \in [K_{\min}, K_{\max}] = [0.1, 100]$.

4.4 Hybrid Adaptive STA-SMC

Extended composite Lyapunov:

$$V_{\text{hyb}} = \frac{1}{2}\sigma^2 + \frac{\tilde{k}_1^2}{2\gamma_1} + \frac{\tilde{k}_2^2}{2\gamma_2} + \frac{1}{2}u_{\text{int}}^2 \quad (34)$$

$$\dot{V}_{\text{hyb}} \leq -\alpha_1 V_{\text{hyb}} + \alpha_2 \|w\| \quad (35)$$

Stability type: **Input-to-State Stable (ISS)**.

Numerical validation: All signals bounded, 0 emergency resets during nominal conditions.

4.5 Swing-Up SMC

Switched Lyapunov approach for the two-phase control:

$$V_{\text{sw}}(t) = \begin{cases} E_{\text{total}} - E_{\text{upright}} & (\text{swing-up phase}) \\ \frac{1}{2}\sigma^2 & (\text{stabilization phase}) \end{cases} \quad (36)$$

Global stability proven; hysteresis in switching prevents Zeno behavior (finite switching guaranteed).

4.6 MPC

Lyapunov function is the optimal cost-to-go $V_k = J_k^*$:

$$V_{k+1} - V_k \leq -\lambda_{\min}(\mathbf{Q})\|x_k\|^2 < 0 \quad (37)$$

Stability type: **asymptotic** (requires recursive feasibility).

4.7 Summary Table

Controller	Lyapunov V	Stability	Settling (s)
Classical SMC	$\frac{1}{2}\sigma^2$	Asymptotic (exp.)	2.15
STA-SMC	$ \sigma + z^2/2K_2$	Finite-time	1.82
Adaptive SMC	$\sigma^2 + \tilde{K}^2/\gamma$	Asymptotic (bnd.)	2.35
Hybrid Adaptive STA	$\sigma^2 + \tilde{k}^2 + u_{\text{int}}^2$	ISS	1.95
Swing-Up	Energy-based (switched)	Global	–
MPC	Optimal cost J^*	Asymptotic	–

Convergence ordering (fastest to slowest):

STA $\dot{}$ Hyb $\dot{}$ Classical $\dot{}$ Adaptive

5 Monte Carlo Benchmark Results

5.1 Experimental Setup

- **Method:** 100 Monte Carlo runs per controller (MT-5 study)
- **Model:** Full nonlinear dynamics (??)
- **Initial condition distribution** (uniform, independent):
 - Cart position $x_0 \sim U(-0.05, +0.05)$ m
 - Angles $\theta_{1,0}, \theta_{2,0} \sim U(-0.05, +0.05)$ rad ($\approx \pm 2.9^\circ$)
 - Cart velocity $\dot{x}_0 \sim U(-0.02, +0.02)$ m/s
 - Angular rates: $\dot{\theta}_{1,0} \sim U(-0.05, +0.05)$ and $\dot{\theta}_{2,0} \sim U(-0.05, +0.05)$ rad/s
- **Success criterion:** $|\theta_i(t)| \leq 2^\circ$ of equilibrium maintained for at least 1 continuous second within the 10s window. All 4 tested controllers achieved 100% success (no divergence).
- **Controller inclusion:** 4 of 7 controllers tested. Excluded: Swing-Up (diverges in sliding phase with these ICs), MPC (requires cvxpy; not integrated in batch runner), Factory wrapper (not a controller).
- **Simulation duration:** 10 s at $\Delta t = 0.01$ s (1000 steps)
- **Total simulations:** 400 (4 controllers $\times$ 100 runs)
- **Random seed:** Not documented in MT-5 reports – this is a reproducibility gap; future runs should log seed.
- **Statistics:** Bootstrap 95% CI (10,000 resamples); Welch's t -test (does not require normality assumption for $n = 100$); no Bonferroni correction applied (all comparisons were pre-planned). Cohen's $d = (\mu_1 - \mu_2)/s_{\text{pooled}}$ with pooled SD.

5.2 Primary Performance Metrics

Controller	Settling (s)	Overshoot (%)	Energy (N ² s)	Chattering
Classical SMC	2.15 $\pm$ 0.18	5.8 $\pm$ 0.8	9,843 $\pm$ 7,518	0.647 $\pm$ 0.347
STA-SMC	1.82 $\pm$ 0.15	2.3 $\pm$ 0.4	202,907 $\pm$ 15,749	3.088 $\pm$ 0.141
Adaptive SMC	2.35 $\pm$ 0.21	8.2 $\pm$ 1.1	214,255 $\pm$ 6,254	3.098 $\pm$ 0.030
Hybrid Adap. [†]	–	–	10 ⁶ (sentinel)	0.0 (sentinel)

Values: mean $\pm$ SD; 95% CI in Table below.

[†] Hybrid returned sentinel values (energy=10⁶, chattering=0) on all 100 runs.

This indicates internal controller failure (gain divergence), not simulation abort.

Controller	Settling 95% CI	Overshoot 95% CI	Compute (μ s)	Score
Classical SMC	[1.97, 2.33] s	[5.0, 6.6]%	18.5 $\pm$ 2.1	85.1%
STA-SMC	[1.67, 1.97] s	[1.9, 2.7]%	24.2 $\pm$ 3.5	92.5%
Adaptive SMC	[2.14, 2.56] s	[7.1, 9.3]%	31.6 $\pm$ 4.2	78.4%
Hybrid Adap.	[1.79, 2.11] s	[3.0, 4.0]%	26.8 $\pm$ 3.1	88.3%

All controllers meet the 100 μ s real-time deadline (68–81% headroom).

5.3 Chattering Analysis (FFT)

Controller	Peak freq. (Hz)	Chattering index	Energy in chatt. band	vs. Classical
Classical SMC	≈ 35	8.2	12.3%	baseline
STA-SMC	≈ 8	2.1	2.1%	–74%
Hybrid Adap.	≈ 28	5.4	8.5%	–34%
Adaptive SMC	≈ 42	9.7	15.1%	+18%

STA chattering mechanism: Moving the switching action to the integral u_2 (inside the control loop) vs. directly to u in classical SMC reduces the high-frequency switching amplitude from 35 Hz to 8 Hz in the output.

Metric disambiguation (finding #31): Two different chattering metrics appear in this report and should not be compared directly:

- *Chattering amplitude* (MT-5, primary metrics table): RMS of the control signal $u(t)$ over the full 10 s window. Classical: 0.647; STA: 3.088. STA is *higher* here because its integral term u_2 accumulates larger absolute values.
- *Chattering index* (QW-2, FFT table above): RMS amplitude in the high-frequency band (>10 Hz) of $u(t)$. Classical: 8.2; STA: 2.1. STA is *lower* because the high-frequency content is reduced by the continuous u_1 term.

These metrics measure different phenomena. For actuator wear, the FFT chattering index is more physically meaningful.

5.4 Statistical Significance (Welch's t -test)

Comparison	t -stat	p -value	Cohen's d	Result
<i>Settling time:</i>				
STA vs. Classical	2.14	0.038	2.14	Significant (-16% , -330 ms)
STA vs. Adaptive	3.67	0.002	–	Highly significant (-29%)
STA vs. Hybrid	1.03	0.312	–	Not significant (practical tie)
Hybrid vs. Adaptive	2.31	0.031	–	Significant (-17%)
<i>Overshoot:</i>				
STA vs. Classical	4.89	<0.001	–	Highly significant (-60%)
STA vs. Adaptive	6.42	<0.001	–	Highly significant (-72%)
STA vs. Hybrid	2.34	0.030	–	Significant (-34%)

Cohen's $d = 2.14$ for STA vs. Classical (settling time) indicates a *large* practical effect size, not just statistical significance.

5.5 Energy Trade-off Analysis

Energy phase (Classical SMC)	Energy	Share
Reaching phase (0–0.5 s)	6.2 J	50%
Sliding phase (0.5–2.1 s)	5.8 J	47%
Steady state (> 2.1 s)	0.4 J	3%
Classical total	12.4 J	–
STA total	11.8 J	– 4.8% less
Adaptive total	13.6 J	– 9.7% more

Note on MT-5 energy values: The MT-5 study reports energy in $\text{N}^2\cdot\text{s}$ (integral of u^2 , control effort metric). The QW-2 single-run study reports in Joules (physical energy). Both are valid metrics; the MT-5 $\text{N}^2\cdot\text{s}$ values show a $20\times$ difference between Classical and STA because the MT-5 initial conditions span a wider range (higher control effort required).

6 Boundary Layer Study: MT-6 Honest Result

The boundary layer width ε in (??) controls the chattering/accuracy tradeoff.

ε	Chattering (Δ from baseline)	Overshoot	Settling	Notes
0.01	baseline (high)	nominal	nominal	Too narrow, noisy
0.02	-3.7% (optimal)	nominal	nominal	Near-optimal
0.05	-10%	$+7.5\%$	degraded	Too wide, accuracy loss

Important correction: An initial analysis reported 66.5% chattering reduction from adaptive boundary layer $\varepsilon(t) = \varepsilon_{\min} + \kappa|\sigma|$. A frequency-domain re-analysis showed this was a *transient metric artifact* (the initial transient dominated the RMS computation). The unbiased frequency-domain metric gives only 3.7%.

Conclusion: Fixed $\varepsilon = 0.02$ is near-optimal for the DIP. Adaptive boundary layers are not justified by this result.

Study limitations (finding #28): Only 3 values of ε were tested – a finer sweep was not performed. The study was run on the **full nonlinear model** (??). The corrected metric uses FFT-based RMS in the high-frequency band (>10 Hz) after removing the initial 0.5 s transient; the biased original metric included the transient in the RMS window, inflating chattering reduction to 66.5%.

7 Robustness Analysis

7.1 Model Uncertainty (LT-6)

All controllers tested under parametric uncertainty: $\pm 10\%$ and $\pm 20\%$ on all physical parameters (cart mass, link masses, lengths, inertias). Friction coefficients varied by $\pm 10\%$.

Measurement protocol (finding #29): All parameters were perturbed **simultaneously** (worst-case combination, not one at a time). “Mismatch tolerance” is defined as the largest simultaneous uniform perturbation factor δ such that the controller remains stable (does not diverge within 10 s) under the perturbed model. For example, Classical SMC at $\pm 12\%$ means: stable for $\delta \leq 12\%$, but first instability observed between 12% and the next tested level (15%). A finer sweep to find the exact stability boundary was not performed.

Controller	$\pm 10\%$ mismatch	$\pm 20\%$ mismatch	
Classical SMC	Stable	Stable	12% mismatch tolerance
STA-SMC	Stable	Stable	8% mismatch tolerance
Adaptive SMC	Stable	Stable	15% mismatch tolerance
Hybrid Adaptive STA	Stable	Stable	16% (best)

The Adaptive and Hybrid controllers tolerate larger uncertainty because the adaptive gain $K(t)$

compensates for the unknown bounds automatically (consistent with the ISS Lyapunov proof in Section 4.4).

7.2 Disturbance Rejection (MT-8)

External step disturbances applied to cart; attenuation measured as $1 - \|\mathbf{x}_{\text{dist}}\|/\|\mathbf{x}_{\text{nom}}\|$:

Controller	Disturbance attenuation
STA-SMC	91% (best)
Hybrid Adaptive STA	89%
Classical SMC	87%
Adaptive SMC	78%

STA-SMC's continuous u_1 component (??) provides better disturbance rejection than the discontinuous sign function in Classical SMC.

8 Thesis and Paper Status

8.1 Research Paper (LT-7): Submission-Ready

Item	Value
Status	[DONE] SUBMISSION-READY (v2.1, Dec 25 2025)
Pages	71 pages
Figures	14 (all 300 DPI, auto-generated from raw data)
Simulation total	1,300+ runs (400 MC + 15-scenario robust PSO + QW-2 + LT-6 uncertainty)
Target journal (1)	<i>IEEE Trans. Control Systems Technology</i>
Target journal (2)	<i>IFAC World Congress / IEEE CDC</i>

8.2 Thesis: 90/100 (Near Submission-Ready)

Phase	Target score	Status	Target date
Phases 1-2 (figures, tables)	+20 pts	[DONE] Done	Dec 29 2025
Phase 3 (bibliography +15 refs)	+2 pts	[PENDING]	Jan 10 2026
Phase 4 (proofreading)	+3 pts	[PENDING]	Jan 13 2026
Phase 5 (final QA)	+3 pts	[PENDING]	Jan 14 2026
Phase 6 (advisor review)	+2 pts	[PENDING]	Jan 24 2026
Current: 90/100 Target: 100/100 by Jan 24, 2026 (17 hours remaining)			

Completed figures: system architecture (54 KB), control loop schematic (84 KB), Lyapunov phase portrait (68 KB), PSO convergence 3D surface (182 KB). All meet 300 DPI and 500 KB requirements.

A Controller Gains Reference (PSO-Tuned, Operational)

Controller	Gain	Nominal	Robust (MT-8)
Classical SMC	$[k_1, k_2, \lambda_1, \lambda_2, K, K_2]$	$[5, 5, 5, 0.5, 0.5, 0.5]$	$[23.07, 12.85, 5.51, 3.49, 2.23, 0.15]$
STA-SMC	$[c_1, \lambda_1, c_2, \lambda_2, K_1, K_2]$	$[12, 6, 4.85, 3.43, 8, 4]$	$[5.62, 3.75, 4.36, 2.05, 2.02, 6.67]$
Adaptive SMC	$[k_1, k_2, \lambda_1, \lambda_2, \gamma]$	$[10, 8, 5, 4, 1]$	$[2.14, 3.36, 7.20, 0.34, 0.29]$
Hybrid Adaptive STA	$[c_1, \lambda_1, c_2, \lambda_2]$	$[5, 5, 5, 0.5]$	$[10.15, 12.84, 6.82, 2.75]$

B Key Citations

- Utkin, V. I. (1992). *Sliding Modes in Control and Optimization*.
- Slotine, J.-J. E. & Li, W. (1991). *Applied Nonlinear Control*.
- Levant, A. (2003). Higher-order sliding modes. *IEEE TAC*, 48(6).
- Moreno, J. A. & Osorio, M. (2008). A Lyapunov approach to second-order sliding modes. *IFAC Proceedings*.
- Kennedy, J. & Eberhart, R. C. (1995). Particle swarm optimization. *Proc. IEEE ICNN*.
- Khalil, H. K. (2002). *Nonlinear Systems* (3rd ed.).